Deep Vora

Boston, MA | 571-356-3000 | deepvora15@gmail.com | LinkedIn | Github | Portfolio

Education

- **Master of Science in Computer Science,** George Mason University — May 2024
  Relevant Coursework: Data Structures & Algorithms, Secure Software Engineering, Database Architecture
  Big Data using Distributed Systems, Data mining using ML, Advance OOPs, Real-Time Embedded Systems, Advance Mathematics

- **Bachelor of Science in Computer Engineering,** University of Mumbai — January 2021
  Relevant Coursework: Data Structures & Algorithms, Cloud Computing, Software Engineering, Distributed Systems,
  Augmented Reality, Computer and Mobile Networking, Data Analysis, Web Development

Skills

Languages: Java, Python, C#, Kotlin, C++, JavaScript, TypeScript, SQL, NOSQL
Databases: MySQL, PostgreSQL, MongoDB, Oracle, Redis, SQL Server (2016)
Frameworks & Tools: Spring Boot, Node.js, Django, React.js, Hibernate, Express.js, Next.js, Angular, Mockito
DevOps & Cloud: Docker, Kubernetes, Terraform, Jenkins, GitLab, CI/CD, Kafka, AWS (EC2, S3, API Gateway, CloudWatch, SQS, Lambda), GCP (Cloud Run, Cloud SQL), Elasticsearch, Azure
Machine Learning: PySpark, Scikit-learn, Hadoop, NLP, Sentiment Analysis, Regression, Apache Spark
Other Skills: Agile Methodologies, REST API Design, Multithreading, Serverless Computing, Automated Testing, Design Patterns

Experience

Open Learning Exchange
Mobile Software Developer — August 2024 - Present
- Designed and implemented high-performance mobile applications using Java and Kotlin, optimizing speed and efficiency for diverse Android devices.
- Led feature development and code optimization, improving user experience and scalability.
- Conducted thorough debugging and code reviews, ensuring robust, maintainable code for Android applications utilizing CouchDB and Gradle.
- Managed the mobile development lifecycle, from architectural decisions to deployment, integrating the latest Android development trends to enhance app performance.

Bitwise Inc
Fullstack Software Engineer — February 2021 - August 2022
- Engineered fault-tolerant microservices with Spring Boot and Resilience4j, reducing defects by 48% and enhancing service reliability
- Developed scalable REST APIs with Java, Spring Boot, and MySQL, integrating them with financial systems and front-end applications in React.js
- Deployed containerized multi-tier applications using Docker and Kubernetes on GCP, implementing CI/CD pipelines with Jenkins and GitLab
- Optimized data streaming by integrating Kafka, ensuring real-time processing of high-volume transactional data
- Automated testing and improved software quality using JUnit, Cypress, and Mockito, reducing production bugs and improving deployment cycles

Publications & Certifications

- AWS Developer Associate — March 2025
- AWS AI Practioner — March 2025
- Self-Driving Car Simulator Research Paper (ISSN 2394 - 7780) — March 2022
- Digital Soil Fertility Analyzer in Chhatrapati Vishwakarma Awards — September 2018

Projects

Receipt Processor Challenge — July 2024
Java, Spring Boot, SQL, AWS (EC2, S3), Kubernetes, Minikube, Docker, React.js, Microservices, JWT
- Designed and deployed a highly available Spring Boot microservice to process receipt data, generating unique transaction IDs and calculating reward points.
- Leveraged AWS ALB with Auto Scaling Groups to ensure efficient traffic distribution and scalability.
- Implemented SSL encryption and integrated Nginx for secure API routing. (Github)

Recommender System using ALS and Collaborative Filtering — May 2024
Python, PySpark, ALS, Collaborative Filtering, Machine Learning, Scikit-learn, Data Science, Distributed applications, Data Modelling
- Model accuracy through hyperparameter tuning.
- Optimized system performance for large-scale datasets, enabling real-time recommendation generation (Github)

Topic Modelling and Sentiment Detection — April 2024
Python, Pandas, PySpark, LDA, pyLDAvis, NLP, Sentiment Analysis, Databricks, Distributed applications, Data Models
- Processed and analyzed large-scale Twitter data using PySpark, extracting meaningful insights through Latent Dirichlet Allocation (LDA) for topic modeling.
- Conducted sentiment analysis to evaluate public opinion trends on political discourse, leveraging pyLDAvis for data visualization. (LDA)

Leadership

- National Service Scheme, *Event Organizer* — August 2022
- Hackathons & Competitions: Finalist in GMU Hackathons, PatriotHacks'23 — October 2023
- Open Source Contributions: Contributor to various open-source projects, including Open Learning Exchange — August 2024